

# Diffusion Models Beat GANs on Image Synthesis

Year: 2021 | Citations: 9979 | Authors: Prafulla Dhariwal, Alex Nichol

---

## Abstract

We show that diffusion models can achieve image sample quality superior to the current state-of-the-art generative models. We achieve this on unconditional image synthesis by finding a better architecture through a series of ablations. For conditional image synthesis, we further improve sample quality with classifier guidance: a simple, compute-efficient method for trading off diversity for fidelity using gradients from a classifier. We achieve an FID of 2.97 on ImageNet  $128 \times 128$ , 4.59 on ImageNet  $256 \times 256$ , and 7.72 on ImageNet  $512 \times 512$ , and we match BigGAN-deep even with as few as 25 forward passes per sample, all while maintaining better coverage of the distribution. Finally, we find that classifier guidance combines well with upsampling diffusion models, further improving FID to 3.94 on ImageNet  $256 \times 256$  and 3.85 on ImageNet  $512 \times 512$ . We release our code at <https://github.com/openai/guided-diffusion>